

Computer Vision

4. Filtering

I-Chen Lin

College of Computer Science,
National Yang Ming Chiao Tung University

Outline

- ▶ Impulse response and convolution.
- ▶ Linear filter and image pyramid.

Textbook:

- David A. Forsyth and Jean Ponce, Computer Vision: A Modern Approach, Prentice Hall, New Jersey, 2003 or 2012.

Some contents are from the reference lecture notes:

- Prof. D. Lowe, Computer Vision, UBC, CA.
- Prof. D.A. Forsyth, Computer Vision, UIUC.
- Prof. J. Rehg, Computer Vision, Georgia Inst. of Tech.
- Prof. T. Darrell, Computer Vision and Applications, MIT.

What's linear filtering?

- ▶ Construct a new image whose pixels are a weighted sum of the original pixel values.
 - ▶ Using the same set of weights at each point.
 - ▶ Mainly based on neighborhood of the target pixels.
- ▶ For example,

6	5	3
4	6	3
1	2	7

Some functions



	5/4	

Linear filtering

- ▶ The simplest and most useful case:
 - ▶ Replace each pixel with a linear combination of its neighbors.
- ▶ The prescription for the linear combination is called the kernel.

6	5	3
4	6	3
1	2	7

 $*$

0	0	0
-1/4	1/2	-1/4
0	0	0

 $=$

	5/4	

kernel

Correlation

I

10	11	10	0	0	1
9	10	11	1	0	1
10	9	10	0	2	1
11	10	9	10	9	11
9	10	11	9	99	11
10	9	9	11	10	10

F

1	1	1
1	1	1
1	1	1

1/9

O


X	X	X	X	X	X
X	10				X
X					X
X					X
X					X
X	X	X	X	X	X

$$1/9.(10 \times 1 + 11 \times 1 + 10 \times 1 + 9 \times 1 + 10 \times 1 + 11 \times 1 + 10 \times 1 + 9 \times 1 + 10 \times 1) = 1/9.(90) = 10$$


Normalized correlation



Zero mean image, -1:1 scale





Zero mean image, -max:max scale




Normalized correlation



Zero mean image, -1:1 scale



Zero mean image, -max:max scale



Finding hands

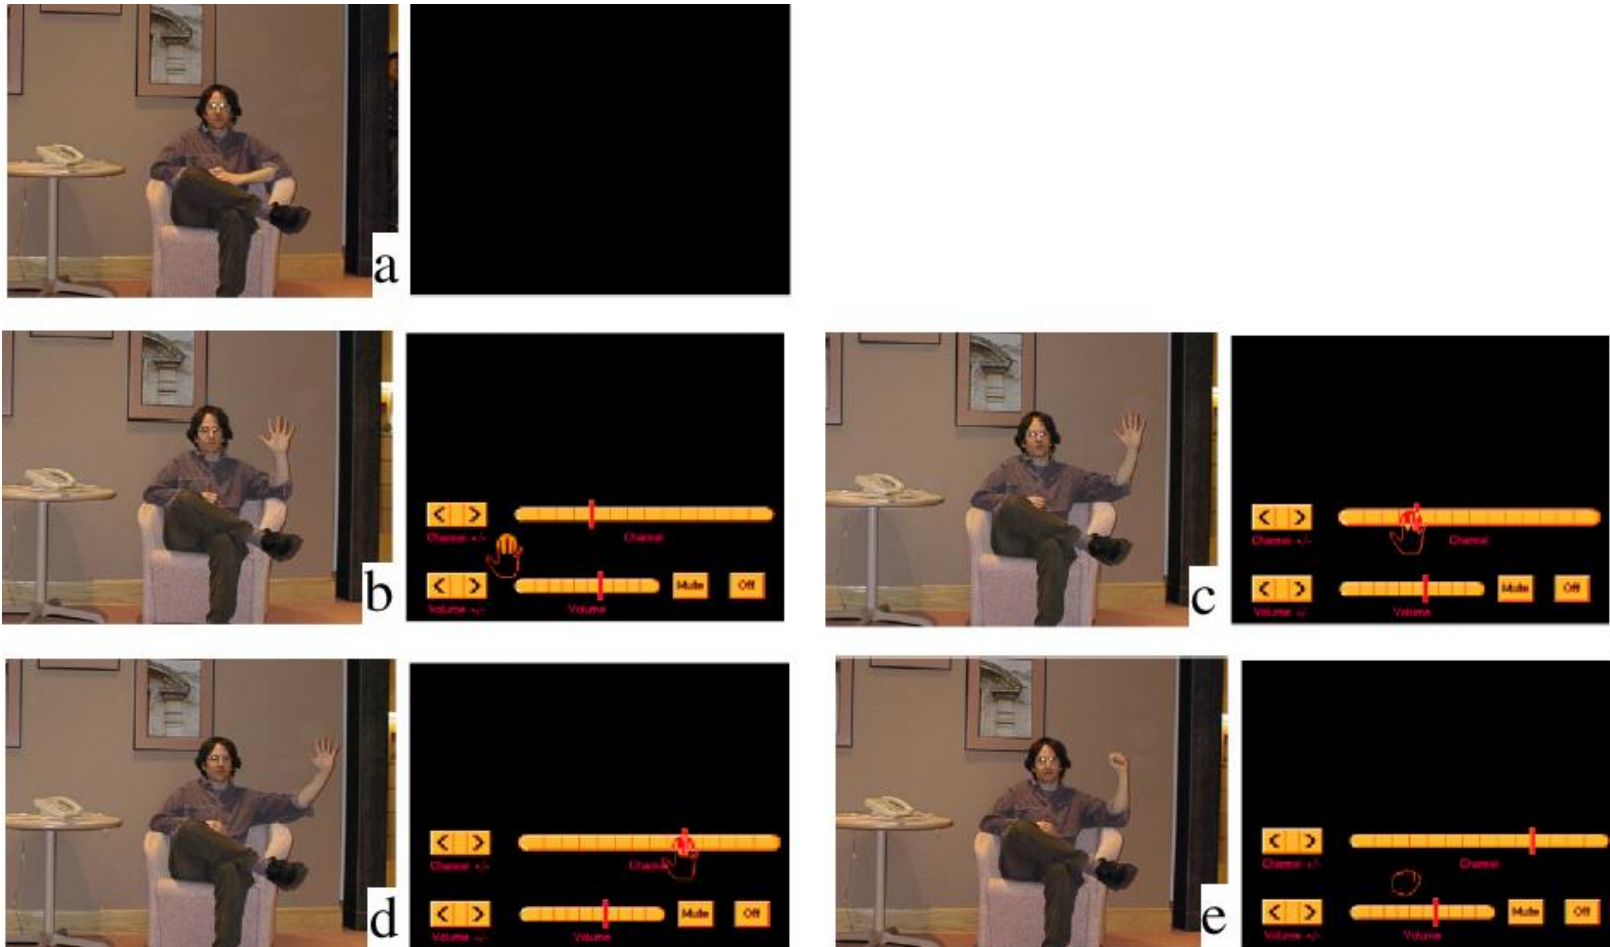


Figure from “Computer Vision for Interactive Computer Graphics,” W.Freeman et al, IEEE Computer Graphics and Applications, 1998.

Convolution

- ▶ The same as correlation with reversed kernels.
- ▶ Represent the linear weights as an image, F , called the kernel.
- ▶ Convolution operation:
 - ▶ Center origin of the kernel F at each pixel location
 - ▶ Multiply weights by corresponding pixels
 - ▶ Set resulting value for each pixel
- ▶ Image, R , resulting from convolution of F with image H , where u, v range over kernel pixels:

$$R_{ij} = \sum_{u,v} H_{i-u,j-v} F_{uv}$$

Warning: the textbook mixes up H and F

Correlation compared to convolution

► Correlation

$$O(x, y) = \sum_{j=-k}^k \sum_{i=-k}^k F(i, j) I(x+i, y+j)$$

► Convolution

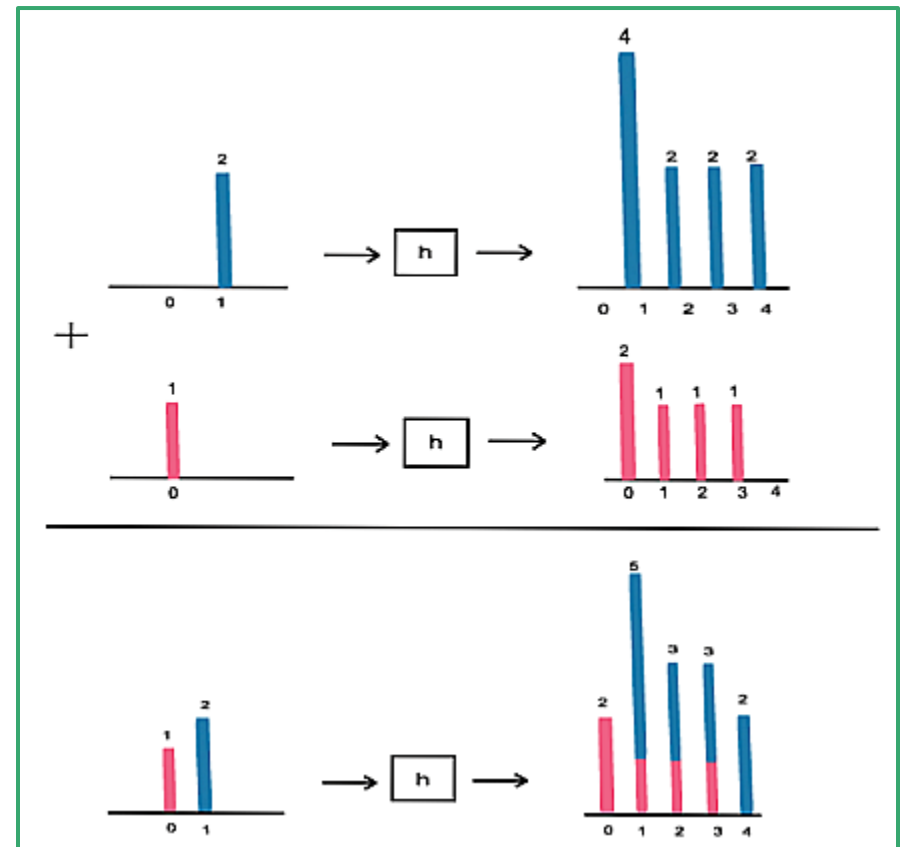

$$\begin{aligned} O(x, y) &= \sum_{j=-k}^k \sum_{i=-k}^k F(i, j) I(x-i, y-j) \\ &= \sum_{j=-k}^k \sum_{i=-k}^k F(-i, -j) I(x+i, y+j) \end{aligned}$$

If $F(i, j) == F(-i, -j)$, then convolution is equivalent to correlation.

Convolution and impulse response

- Convolution can be viewed as a sequence of impulse responses.

$$O(x) = \sum_{i=-k}^k h(i)I(x-i)$$



Examples



Original

0	0	0
0	1	0
0	0	0




Filtered (no change)




Original

0	0	0
0	0	1
0	0	0




Shifted right By 1 pixel

Examples (cont.)



Original

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



Blur (with a box filter)



Original



$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$





Sharpening filter

Sharpening

- ▶ Sharpening filter
 - ▶ Accentuates differences with local average




Examples: Smoothing with a box filter



Smoothing with a Gaussian

- ▶ Averaging does not model defocussed lens well
- ▶ Impulse response should be fuzzy blob, but a box filter would give a little square.
- ▶ Gaussian gives a better model of a fuzzy blob.





Gaussian kernel

- ▶ Weight contributions of neighboring pixels by nearness.

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



- ▶ For efficiency, usually use integers as weights, e.g.: $0.125 *$
 - ▶ Don't forget normalizing the weight sum to 1.

	1	
1	4	1
	1	



0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003


Smoothing with a Gaussian



Additive Gaussian noise



$\sigma = 1$




$\sigma = 16$

Noise and smoothing

- ▶ Smoothing reduces pixel noise:




- ▶ Each row shows smoothing with Gaussians of different width

- ▶ Each column shows different amounts of Gaussian noise.



Noise and smoothing

- ▶ Filtered noise is sometimes useful
 - ▶ E.g. in textures synthesis.



Efficient implementation

- ▶ Both the box filter and the Gaussian filter are separable into two 1D convolutions:
 - ▶ First convolve each row with a 1D filter
 - ▶ Then convolve each column with a 1D filter.

$$\begin{aligned} G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

- ▶ A 2D Gaussian can be expressed as the product of Gaussian of X and Y .

Other noise and smoothing



Could be salt-and-pepper noise.

Figures from lecture notes D.A. Forsyth, Computer Vision, UIUC.

Differentiation

- Recall, for 2D function, $f(x, y)$:

$$\frac{\partial f}{\partial x} = \lim_{\varepsilon \rightarrow 0} \left(\frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon} \right)$$



- This is linear and shift invariant, so must be the result of a convolution.
- We could approximate this as symmetric finite difference:

$$\frac{\partial h}{\partial x} \approx \frac{h_{i+1,j} - h_{i-1,j}}{2} \quad H = \frac{1}{2} \begin{bmatrix} 0 & 0 & 0 \\ 1 & 0 & -1 \\ 0 & 0 & 0 \end{bmatrix}$$


- Issue: noise
 - Smooth before differentiation
 - Two convolutions to smooth, then differentiate?

Differentiation


- ▶ Because differentiation is convolution, and convolution is associative.
 - ▶ $((f ** g) ** h) = (f ** (g ** h))$
- ▶ We can use a derivative of Gaussian filter.




Scale affects derivatives



1 pixel



3 pixel




7 pixel

Aliasing


- ▶ Can't shrink an image by taking every second pixel
- ▶ If we do so, characteristic errors appear
 - ▶ Typically, small phenomena look bigger; fast phenomena can look slower
- ▶ Common phenomenon
 - ▶ Wagon wheels rolling the wrong way in movies
 - ▶ Checkerboards misrepresented in ray tracing
 - ▶ Striped shirts look funny on color television

Re-sampling

- Resample the checkerboard by taking one sample at each circle.




Simple sampling, but ...



Smoothing as low-pass filtering


- ▶ High frequencies lead to trouble with sampling.
- ▶ Suppress high frequencies before sampling !
- ▶ Truncate high frequencies in FT or convolve with a low-pass filter.



The Gaussian pyramid


- ▶ Create each level from previous one:
 - ▶ smooth and sample
- ▶ Smooth with Gaussians, in part because
 - ▶ a Gaussian * Gaussian = another Gaussian
 - ▶ $G(x) * G(y) = G(\sqrt{x^2 + y^2})$
- ▶ Gaussians are low pass filters, so the representation is redundant once smoothing has been performed

The Gaussian pyramid



1D Fourier Transform

- ▶ Represent functions on a new basis .
- ▶ Think of functions as vectors, with many components.



2D Fourier Transform

- ▶ Represent function on a new basis .
 - ▶ Think of functions as vectors, with many components.
- ▶ We now apply a linear transformation to transform the basis dot product with each basis element
- ▶ Basis elements have the form

$$e^{-i\theta} = \cos \theta + i \sin \theta$$

$$e^{-i2\pi(ux+vy)}$$

$$F(f(x,y))(u,v) = \iint_{\mathbb{R}^2} f(x,y) e^{-i2\pi(ux+vy)} dx dy$$

Discrete Fourier Transform

discrete domain

Forward transform

$$F[m, n] = \sum_{k=0}^{M-1} \sum_{l=0}^{N-1} f[k, l] e^{-2\pi i \left(\frac{km}{M} + \frac{ln}{N} \right)}$$

Inverse transform

$$f[k, l] = \frac{1}{MN} \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} F[m, n] e^{2\pi i \left(\frac{km}{M} + \frac{ln}{N} \right)}$$

Real part of basis element

- The magnitude of the vector (u, v) gives a frequency
- Its direction gives an orientation






Phase and magnitude

- ▶ $f(x,y)$ is the image and is REAL
- ▶ $F(u,v)$ (abbreviate as F) is in general, COMPLEX.
- ▶ Fourier transform is complex and difficult for visualization.
- ▶ Instead, think of phase and magnitude
 - ▶ Magnitude = magnitude of complex transform
$$= \text{sqrt}(\text{REAL}(F)^2 + \text{IMAGINARY}(F)^2)$$
 - ▶ Phase = phase of the complex transform
$$= \text{atan} (\text{IMAGINARY}(F) / \text{REAL}(F))$$


Phase and magnitude



Image $f(x,y)$



Magnitude of $F(u,v)$



Phase of $F(u,v)$

Phase and magnitude

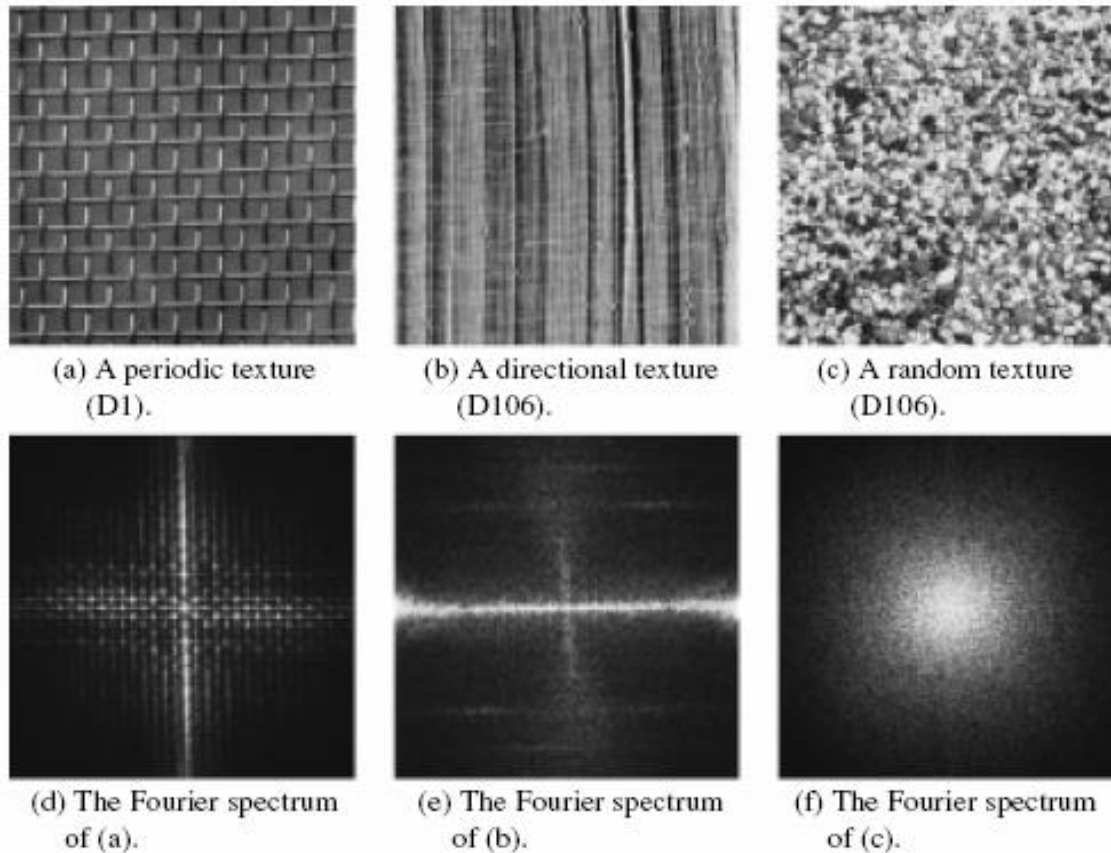



Figure from Lee and Chen, A New Method for Coarse Classification of Textures and Class Weight Estimation for Texture Retrieval, Pattern Recognition and Image Analysis, Vol. 12, No. 4, 2002, pp. 400–410.

Phase and magnitude

- Reconstruct with one's phase, and magnitude of the other.




Reconstruct with zebra phase, cheetah magnitude



Reconstruct with cheetah phase, zebra magnitude

The process and application of FFT



Source: www.mediacy.com/apps/fft.htm, Image Pro Plus FFT Example.
Slides from Earl F. Glynn, "Fourier Analysis and Image Processing"